


CLAUDE BEYOND PROMPTING

Today we will learn how go
beyond prompting with Claude

What Claude actually “sees”

When you open Claude Code in a project directory, it builds its context from four sources — in this order. Everything it does in a session is shaped by this stack.



No magic — just text it can see.

1

System prompt

Built into Claude Code — you don't control this

2

CLAUDE.md

Your project-level instructions — read first, every session

3

Conversation history

Everything you've said and Claude has replied

4

Tool results

File contents, command output, search results

CLAUDE.md — your project's briefing document

The instructions you'd give a new teammate on day one. Claude reads it automatically at the start of every session — no pasting into chat.



What are we building?

The project's purpose and shape



What are the major rules?

Conventions, style, structure to follow



What should Claude NOT do?

Hard constraints and no-go zones



KEY INSIGHT

**A good CLAUDE.md prevents 80%
of “Claude did something weird”
moments.**

If you don't give it constraints, it will invent its own.

Two files, two jobs



CLAUDE.md

- Project-level rules
- Applies to every session
- The broader understanding of the project



task.md

- One specific piece of work
- Scoped to the session at hand
- What to do right now

you > Read task.md and do what it says

How you hand the task over

Why this matters

Without context files



You repeat yourself every single session



Claude makes assumptions you didn't want



More time correcting than building

With context files



Claude starts already knowing the rules



You hand it a task and let it work



Instructions are version-controlled with your code

Demo

- We pick a goal
- Write Claude.md and task.md files for the goal
- Understand context managed within Claude sessions
- Check deliverables